

# SPRINT PLAN — Evidence Ingestion, Traceability & Attestation Workflows

**Created:** 2026-02-06 18:25 MST **Baseline commit:** 32e502e (157/157 tests passing) **Governing doc:** docs/architecture/REGULATORY\_EXECUTION\_PLATFORM\_ARCHITECTURE.md **Priority order:** Regulatory Twin → Traceability → Gap Detection → Dashboard

---

## CURRENT STATE (as of commit 32e502e)

### What EXISTS in the database

| Table           | Status          | Notes                                  |
|-----------------|-----------------|----------------------------------------|
| organizations   | ✓ Created       | Base schema                            |
| users           | ✓ Created       | Minimal — needs role/profile expansion |
| products        | ✓ Created       | Has org FK                             |
| device_versions | ✓ Created       | Has regulatory_twin_json JSONB         |
| org_members     | ✓ Created       | Multi-org membership                   |
| ai_runs         | ✓ Created + RLS | Full provenance schema                 |
| trace_links     | ✓ Created + RLS | Polymorphic link table                 |
| artifacts       | ✓ Created + RLS | Evidence object registry               |
| artifact_links  | ✓ Created + RLS | Evidence-to-target links               |
| attestations    | ✓ Created + RLS | Human sign-off records                 |

### What DOES NOT EXIST yet

| Entity (from architecture) | DB Table | Python Model | API | Tests |
|----------------------------|----------|--------------|-----|-------|
| Intended use               | ✗        | ✗            | ✗   | ✗     |
| Claims                     | ✗        | ✗            | ✗   | ✗     |
| Hazards                    | ✗        | ✗            | ✗   | ✗     |

| Entity (from architecture) | DB Table | Python Model | API | Tests |
|----------------------------|----------|--------------|-----|-------|
| Harms                      | ✗        | ✗            | ✗   | ✗     |
| Risk controls              | ✗        | ✗            | ✗   | ✗     |
| Verification tests         | ✗        | ✗            | ✗   | ✗     |
| Validation tests           | ✗        | ✗            | ✗   | ✗     |
| Evidence items             | ✗        | ✗            | ✗   | ✗     |
| Labeling assets            | ✗        | ✗            | ✗   | ✗     |
| Submission targets         | ✗        | ✗            | ✗   | ✗     |

### What EXISTS in Python (but not connected to DB)

| Component                 | Status                           |
|---------------------------|----------------------------------|
| DeviceInfo Pydantic model | ✓ In-memory only, no persistence |
| ClassificationEngine      | ✓ Works, no DB storage           |
| PathwayEngine             | ✓ Works, no DB storage           |
| ai_runs_logger.py         | ✓ Writes to ai_runs table        |
| supabase_client.py        | ✓ Connection helper              |

## SPRINT STRUCTURE

Three sprints, each ~1 week. Each sprint delivers:

- SQL migration(s)
- Python models + persistence layer
- Tests (unit + integration)
- Snapshot update

### Engineering principles (from governing architecture):

- Structure first, AI second

- Additive changes only
  - Version everything
  - No regulatory hallucinations
  - Human-in-the-loop mandatory
- 

## SPRINT 1 — REGULATORY TWIN CORE ENTITIES

**Goal:** Create the structured data foundation that makes everything else possible. Every device becomes structured data — not just a Pydantic model in memory.

### Deliverable 1A: Core Domain Tables Migration

**File:** `scripts/migrations/2026-02-07_regulatory_twin_core.sql`

New tables (all org-scoped, all with RLS):

### intended\_uses

- id (uuid PK)
- organization\_id (FK → organizations)
- device\_version\_id (FK → device\_versions)
- statement (text NOT NULL) -- the intended use statement
- indications (jsonb DEFAULT '[]') -- list of indications
- contraindications (jsonb DEFAULT '[]')
- target\_population (text)
- use\_environment (text) -- clinical, home, point-of-care
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → intended\_uses, self-ref for versioning)
- created\_at (timestamptz)

### claims

- id (uuid PK)
- organization\_id (FK → organizations)
- device\_version\_id (FK → device\_versions)
- claim\_type (text NOT NULL) -- safety, performance, usability
- statement (text NOT NULL)
- evidence\_basis (text) -- clinical, bench, lit review
- status (text DEFAULT 'draft') -- draft, under\_review, accepted, rejected
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → claims)
- created\_at (timestamptz)

### hazards

- id (uuid PK)
- organization\_id (FK → organizations)
- device\_version\_id (FK → device\_versions)
- hazard\_category (text NOT NULL) -- electrical, biological, software, use-error
- description (text NOT NULL)
- foreseeable\_sequence (text) -- how hazard leads to harm
- severity (text) -- negligible, marginal, critical, catastrophic
- probability (text) -- improbable, remote, occasional, probable, frequent
- risk\_level\_pre (text) -- pre-mitigation: low, medium, high, unacceptable
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → hazards)
- created\_at (timestamptz)

### harms

- id (uuid PK)
- organization\_id (FK → organizations)
- hazard\_id (FK → hazards)
- harm\_type (text NOT NULL) -- injury, death, misdiagnosis, delay
- description (text NOT NULL)
- severity (text NOT NULL)
- affected\_population (text)
- created\_by (FK → users)
- created\_at (timestamptz)

#### risk\_controls

- id (uuid PK)
- organization\_id (FK → organizations)
- hazard\_id (FK → hazards)
- control\_type (text NOT NULL) -- design, protective, information
- description (text NOT NULL)
- risk\_level\_post (text) -- post-mitigation residual risk
- implementation\_status (text DEFAULT 'planned') -- planned, implemented, verified
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → risk\_controls)
- created\_at (timestamptz)

#### verification\_tests

- id (uuid PK)
- organization\_id (FK → organizations)
- device\_version\_id (FK → device\_versions)
- test\_type (text NOT NULL) -- bench, software, biocompat, electrical
- title (text NOT NULL)
- protocol\_ref (text) -- reference to test protocol doc
- acceptance\_criteria (text NOT NULL)
- result\_summary (text)
- pass\_fail (text) -- pass, fail, conditional, pending
- tested\_at (timestamptz)
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → verification\_tests)
- created\_at (timestamptz)

#### validation\_tests

- id (uuid PK)
- organization\_id (FK → organizations)
- device\_version\_id (FK → device\_versions)
- test\_type (text NOT NULL) -- usability, clinical, simulated\_use

- title (text NOT NULL)
- protocol\_ref (text)
- acceptance\_criteria (text NOT NULL)
- result\_summary (text)
- pass\_fail (text)
- participant\_count (integer)
- tested\_at (timestamptz)
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → validation\_tests)
- created\_at (timestamptz)

#### evidence\_items

- id (uuid PK)
- organization\_id (FK → organizations)
- device\_version\_id (FK → device\_versions)
- evidence\_type (text NOT NULL) -- test\_report, lit\_review, clinical\_data, standard\_ref
- title (text NOT NULL)
- description (text)
- source\_ref (text) -- external reference (standard number, paper DOI)
- artifact\_id (FK → artifacts) -- link to stored file/content
- strength (text) -- strong, moderate, weak, insufficient
- status (text DEFAULT 'draft')
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → evidence\_items)
- created\_at (timestamptz)

#### labeling\_assets

- id (uuid PK)
- organization\_id (FK → organizations)
- device\_version\_id (FK → device\_versions)
- asset\_type (text NOT NULL) -- ifu, label, packaging, e-labeling
- title (text NOT NULL)
- content\_ref (text) -- storage URI or artifact link
- language (text DEFAULT 'en')
- regulatory\_market (text) -- CA, US, EU
- artifact\_id (FK → artifacts)
- status (text DEFAULT 'draft')
- created\_by (FK → users)
- version (integer NOT NULL DEFAULT 1)
- supersedes\_id (FK → labeling\_assets)
- created\_at (timestamptz)

```
submission_targets
- id (uuid PK)
- organization_id (FK → organizations)
- device_version_id (FK → device_versions)
- regulatory_body (text NOT NULL) -- health_canada, fda, eu_mdr
- submission_type (text NOT NULL) -- mdl, 510k, de_novo, pma, ce_mark
- target_date (date)
- status (text DEFAULT 'planning') -- planning, preparing, submitted, approved, rejected
- reference_number (text) -- assigned by regulator
- created_by (FK → users)
- created_at (timestamptz)
```

## Key design decisions:

- Every table has `organization_id` for RLS consistency
- Every mutable table has `version` + `supersedes_id` for immutable versioning (Law 7)
- All use `created_by` FK to users for audit trail
- Hazard→Harm→Risk\_Control chain is explicit (not polymorphic) for safety
- Evidence\_items link to artifacts (stored files) via FK
- Status fields use text enums (not Postgres enums) for flexibility

## Deliverable 1B: Pydantic Models

**File:** `src/core/regulatory_twin.py`

Pydantic models mirroring every table above. Each model includes:

- Validators for status/type fields
- `.to_db_dict()` method for persistence
- `@classmethod from_db_row()` for reading
- Proper Optional fields for nullable columns

## Deliverable 1C: Persistence Layer

**File:** `src/persistence/twin_repository.py`

CRUD operations for each entity:

- `create_intended_use()`, `get_intended_uses_for_device()`, etc.
- All operations scoped to organization
- Uses existing `supabase_client.py` pattern
- Best-effort (never crashes app on DB failure, matching ai\_runs\_logger pattern)

## **Deliverable 1D: Tests**

### **Files:**

- `tests/unit/test_regulatory_twin_models.py` — Pydantic validation (40+ tests)
- `tests/unit/test_regulatory_twin_migration.py` — SQL file checks (30+ tests, same pattern as RLS tests)
- `tests/integration/test_twin_persistence.py` — DB round-trip tests

## **Sprint 1 Exit Criteria**

- Migration creates all 10 new tables
  - RLS enabled on all 10 new tables
  - All tables have version + supersedes\_id columns
  - Pydantic models for all 10 entities
  - Persistence CRUD for at least: intended\_uses, claims, hazards, risk\_controls, evidence\_items
  - 200+ total tests passing
  - Snapshot updated
  - Committed to main
- 

## **SPRINT 2 — TRACEABILITY ENGINE + EVIDENCE INGESTION**

**Goal:** Make trace\_links operational. Connect the dots: claim → hazard → risk\_control → verification → evidence. This is the PRIMARY WEDGE per governing architecture.

## **Deliverable 2A: Traceability Service**

**File:** `src/core/traceability.py`

The brain of the system. Implements:

```
python
```

```
class TraceabilityEngine:
```

```
    """
```

Creates and queries regulatory trace links.

Supports the full chain:

claim → mitigated\_by → risk\_control

risk\_control → verified\_by → verification\_test

verification\_test → supported\_by → evidence\_item

hazard → causes → harm

hazard → mitigated\_by → risk\_control

claim → supported\_by → evidence\_item

```
"""
```

```
VALID_RELATIONSHIPS = {
```

(“claim”, “hazard”): [“addresses”],

(“claim”, “evidence\_item”): [“supported\_by”],

(“hazard”, “harm”): [“causes”, “may\_cause”],

(“hazard”, “risk\_control”): [“mitigated\_by”],

(“risk\_control”, “verification\_test”): [“verified\_by”],

(“risk\_control”, “validation\_test”): [“validated\_by”],

(“verification\_test”, “evidence\_item”): [“supported\_by”],

(“validation\_test”, “evidence\_item”): [“supported\_by”],

(“evidence\_item”, “artifact”): [“documented\_in”],

```
}
```

```
def create_link(self, source_type, source_id, target_type, target_id,  
               relationship, rationale, created_by) -> TraceLink
```

```
def get_links_from(self, source_type, source_id) -> List[TraceLink]
```

```
def get_links_to(self, target_type, target_id) -> List[TraceLink]
```

```
def get_full_chain(self, claim_id) -> TraceChain
```

““Follow all links from a claim down to evidence.””

```
def get_coverage_report(self, device_version_id) -> CoverageReport
```

““For each claim, show: linked hazards, controls, tests, evidence.””

```
def validate_link(self, source_type, target_type, relationship) -> bool
```

““Check if the relationship is valid per VALID\_RELATIONSHIPS.””

## **Deliverable 2B: Evidence Ingestion Service**

**File:** `src/core/evidence_ingestion.py`

```
python
```

```
class EvidenceIngestionService:
```

```
    """
```

```
    Ingest evidence and connect it to the regulatory twin.
```

Workflow:

1. Create artifact (file metadata + content hash)
2. Create evidence\_item (typed, with strength assessment)
3. Create trace\_link to the relevant claim/test/control
4. Optionally log AI assist via ai\_runs if AI helped classify

```
    """
```

```
def ingest_evidence(self, device_version_id, evidence_type, title,  
                    artifact_data, linked_to) -> EvidenceItem
```

```
def bulk_ingest(self, device_version_id, items: List[dict]) -> List[EvidenceItem]
```

```
def get_evidence_for_claim(self, claim_id) -> List[EvidenceItem]
```

```
def get_unlinked_evidence(self, device_version_id) -> List[EvidenceItem]
```

```
    """Find evidence items not connected to any claim/test/control."""
```

## **Deliverable 2C: Attestation Workflow Service**

**File:** `src/core/attestation_service.py`

```
python
```

```

class AttestationService:
    """
    Human-in-the-loop sign-off on artifacts and links.

    Every AI output, every trace link, every evidence assessment
    can be attested to by a human reviewer.
    """

    ATTESTATION_TYPES = [
        "reviewed",      # human reviewed the content
        "approved",      # human approved for regulatory use
        "rejected",      # human rejected, needs rework
        "acknowledged",  # human saw it, no opinion
    ]

    def attest_artifact(self, artifact_id, attested_by, attestation_type, note)

    def attest_link(self, artifact_link_id, attested_by, attestation_type, note)

    def get_attestation_status(self, artifact_id) -> AttestationStatus

    def get_unattested_items(self, organization_id) -> List[dict]
        """Find artifacts and links that have not been reviewed."""

    def get_attestation_audit_trail(self, artifact_id) -> List[Attestation]

```

## Deliverable 2D: API Endpoints

**File:** `src/api/traceability_routes.py`

New endpoints (added to existing FastAPI app):

```
POST /api/v1/trace-links — create a trace link  
GET /api/v1/trace-links/{id} — get a trace link  
GET /api/v1/trace-chains/{claim_id} — get full chain from claim to evidence  
GET /api/v1/coverage/{device_version_id} — coverage report
```

```
POST /api/v1/evidence — ingest evidence item  
GET /api/v1/evidence/{device_version_id} — list evidence for device  
GET /api/v1/evidence/unlinked/{device_version_id} — unlinked evidence
```

```
POST /api/v1/attestations — create attestation  
GET /api/v1/attestations/pending/{org_id} — unattested items  
GET /api/v1/attestations/trail/{artifact_id} — audit trail
```

## Deliverable 2E: Tests

### Files:

- `tests/unit/test_traceability.py` — link validation, chain traversal (40+ tests)
- `tests/unit/test_evidence_ingestion.py` — ingest, bulk, unlinked detection (25+ tests)
- `tests/unit/test_attestation_service.py` — attest, reject, audit trail (25+ tests)
- `tests/integration/test_trace_chain_flow.py` — full claim→evidence flow (10+ tests)
- `tests/api/test_traceability_endpoints.py` — API route tests (20+ tests)

## Sprint 2 Exit Criteria

- TraceabilityEngine creates/queries links with validation
- Full chain traversal: claim → hazard → control → test → evidence
- Evidence ingestion with artifact creation
- Attestation workflow (attest, reject, audit trail)
- API endpoints for all three services
- Coverage report per device version
- Unlinked evidence detection
- 320+ total tests passing
- Snapshot updated
- Committed to main

---

## SPRINT 3 — GAP DETECTION ENGINE + READINESS ASSESSMENT

**Goal:** The highest-value feature. Rules that surface what's missing, weak, or inconsistent. Per architecture:

"Never say 'You are submission ready.' Always say 'Readiness assessment based on configured expectations.'"

## Deliverable 3A: Gap Detection Rules Engine

File: [src/core/gap\\_engine.py](#)

```
python

class GapDetectionEngine:
    """
    Rules-based engine that evaluates regulatory readiness.

    Rules are versioned, deterministic, and explainable.
    Each rule produces a GapFinding with severity, description, and remediation.
    """

    def evaluate(self, device_version_id) -> GapReport:
        """Run all rules against a device version."""

    def evaluate_rule(self, rule_id, device_version_id) -> List[GapFinding]:
        """Run a single rule."""

    def get_rules(self) -> List[GapRule]:
        """List all active rules with descriptions."""

class GapRule:
    """A single detection rule. Deterministic. Explainable."""
    id: str
    name: str
    description: str
    severity: str      # critical, major, minor, info
    category: str     # coverage, completeness, consistency, evidence_strength
    version: int
    evaluate: Callable # returns List[GapFinding]
```

Initial rule set (Health Canada focus):

| Rule ID | Name                                | What it checks                                          | Severity |
|---------|-------------------------------------|---------------------------------------------------------|----------|
| GAP-001 | Unmitigated hazards                 | Hazards with no linked risk_control                     | CRITICAL |
| GAP-002 | Unverified controls                 | Risk controls with no linked verification_test          | CRITICAL |
| GAP-003 | Unsupported claims                  | Claims with no linked evidence_item                     | MAJOR    |
| GAP-004 | Missing intended use                | Device version with no intended_use record              | CRITICAL |
| GAP-005 | Weak evidence                       | Evidence items with strength = 'weak' or 'insufficient' | MAJOR    |
| GAP-006 | Untested claims                     | Claims with no linked verification OR validation test   | MAJOR    |
| GAP-007 | No submission target                | Device version with no submission_target                | MINOR    |
| GAP-008 | Unattested AI outputs               | ai_runs linked to artifacts but not attested            | MAJOR    |
| GAP-009 | Missing labeling                    | Device version with no labeling_assets                  | MAJOR    |
| GAP-010 | Incomplete risk chain               | Hazard → harm → control chain has breaks                | CRITICAL |
| GAP-011 | Draft evidence only                 | All evidence_items in 'draft' status                    | MAJOR    |
| GAP-012 | No clinical evidence (Class III/IV) | Class III/IV with no clinical evidence type             | CRITICAL |

### Deliverable 3B: Readiness Assessment

File: `src/core/readiness.py`

```
python
```

```

class ReadinessAssessment:
    """
    Aggregates gap findings into a readiness score.

    NEVER says "compliant" or "ready."
    ALWAYS says "Readiness assessment based on configured expectations."
    """

    def assess(self, device_version_id) -> ReadinessReport:
        """
        Returns:
        - overall_readiness_score: float (0.0 - 1.0)
        - category_scores: dict[str, float]
        - gap_findings: List[GapFinding]
        - critical_blockers: List[GapFinding]
        - summary: str (regulatory-safe language)
        """

    def generate_summary(self, report: ReadinessReport) -> str:
        """
        Generates human-readable summary.
        Uses ONLY approved regulatory language.
        """

```

## Deliverable 3C: API Endpoints

```

GET /api/v1/gaps/{device_version_id} — full gap report
GET /api/v1/gaps/{device_version_id}/critical — critical gaps only
GET /api/v1/readiness/{device_version_id} — readiness assessment
GET /api/v1/rules — list all gap rules

```

## Deliverable 3D: Tests

### Files:

- `tests/unit/test_gap_engine.py` — each rule individually tested (50+ tests)
- `tests/unit/test_readiness.py` — scoring, language safety (20+ tests)
- `tests/regulatory/test_gap_rules.py` — regulatory correctness of rules (15+ tests)
- `tests/integration/test_gap_detection_flow.py` — end-to-end (10+ tests)
- `tests/api/test_gap_endpoints.py` — API route tests (10+ tests)

## Sprint 3 Exit Criteria

- 12 gap detection rules implemented and tested
  - Each rule produces explainable findings with severity
  - Readiness assessment with category scores
  - ALL output uses regulatory-safe language (no "compliant", no "ready")
  - API endpoints for gaps, readiness, rules
  - 430+ total tests passing
  - Snapshot updated
  - Committed to main
- 

## CUMULATIVE MILESTONE TRACKER

| Metric             | Baseline (now) | Sprint 1 | Sprint 2 | Sprint 3 |
|--------------------|----------------|----------|----------|----------|
| DB tables          | 9              | 19       | 19       | 19       |
| RLS-enabled tables | 9              | 19       | 19       | 19       |
| Python models      | 6              | 16       | 16       | 18       |
| Service classes    | 2              | 3        | 6        | 8        |
| API endpoints      | 6              | 6        | 15       | 19       |
| Total tests        | 157            | 230+     | 320+     | 430+     |
| Trace link types   | 0              | 0        | 9        | 9        |
| Gap rules          | 0              | 0        | 0        | 12       |

---

## TECH DEBT TO ADDRESS ALONGSIDE

| Item                                | Priority | When                                            |
|-------------------------------------|----------|-------------------------------------------------|
| 34 mypy errors                      | MEDIUM   | Fix progressively per sprint                    |
| public.users vs auth.users decision | HIGH     | Before Sprint 1 migration                       |
| Supabase cloud deployment           | HIGH     | After Sprint 2 (need real auth for RLS testing) |

| Item                        | Priority | When                                   |
|-----------------------------|----------|----------------------------------------|
| S3 document storage         | MEDIUM   | Sprint 2 (evidence needs file storage) |
| Pre-commit mypy enforcement | LOW      | After mypy errors resolved             |

## RISK REGISTER

| Risk                                          | Impact | Mitigation                                                      |
|-----------------------------------------------|--------|-----------------------------------------------------------------|
| Schema changes break existing tests           | HIGH   | All migrations idempotent, test existing 157 before each sprint |
| Supabase auth.users vs public.users confusion | HIGH   | Decide BEFORE Sprint 1, document in CLAUDE.md                   |
| Scope creep on Regulatory Twin entities       | MEDIUM | Stick to architecture entities only, no extras                  |
| Gap rules too rigid                           | LOW    | Version rules, allow disable/enable per org                     |
| No real multi-user testing                    | HIGH   | Supabase cloud deployment before Sprint 3                       |

## WHAT THIS PLAN DOES NOT INCLUDE (deferred)

These are Phase 4-5 per architecture, deliberately excluded:

- Document Orchestration (generating submission documents from structured data)
- Deficiency Response Copilot
- Regulatory Knowledge Graph construction
- Real UI/dashboard (Streamlit or React)
- AI-assisted link recommendations
- S3 signed URL document storage

**Reason:** The architecture says "Earn complexity." We build the data foundation and deterministic logic first. AI and UI come after structure is solid.

*End of sprint plan*